

Pinchy Platform Overview

Pinchy is an enterprise AI agent platform designed for organizations that need to run autonomous AI agents with proper governance and control. Built on top of the OpenClaw runtime, Pinchy adds the enterprise layer that businesses require.

Key features include role-based access control, comprehensive audit trails, and a plugin architecture that allows administrators to precisely control what each agent can and cannot do. Every action is logged and traceable.

Architecture

The platform uses a layered architecture where OpenClaw provides the agent runtime capabilities while Pinchy wraps these with permission checks, user management, and audit logging. This separation ensures that Pinchy benefits from OpenClaw improvements without forking the runtime.

Deployment Guide

Pinchy is deployed using Docker Compose, which orchestrates all required services including the web application, PostgreSQL database, and OpenClaw runtime. A single docker compose up command brings the entire stack online.

The self-hosted deployment model ensures that all data remains within the organization network. No external API calls are required when using local models through Ollama integration. This makes Pinchy suitable for air-gapped environments.

DRAFT - Internal Use Only

Roadmap Notes - Q3 2025

- Plugin marketplace beta launch
- Cross-channel workflow engine
- Advanced RBAC with team scoping
- SSO integration (SAML, OIDC)

These items are subject to change.